

# CISD 43 Final Project

June 9, 2024

## 0.1 Title here ( i.e., House Price Prediction)

```
[1]: # Edit all the Mardown cells below with the appropriate information
# Run all cells, containing your code
# Save this Jupyter with the outputs of your executed cells
# PS: Save again the notebook with this outcome.
# PSPS: Don't forget to include the dataset in your submission
```

**Team:** \* Riker Wachtler

**Course:** CISD 43 – BIG DATA (Spring, 2024)

### 0.1.1 Problem Statement

- This project is about house price predictions. I wasn't very creative, so I just used the sample suggestion of houses prices.
- **Keywords:** House price prediction, real estate

### 0.1.2 Required packages

- If you are using Jupyter Notebook via Anaconda with **pip** properly configured, the following lines should successfully install the required modules. They are commented out by default as to not attempt to re-install if you were to fully re-run the notebook

```
[2]: #!pip install pandas
#!pip install numpy
#!pip install matplotlib
#!pip install seaborn
#!pip install sklearn
#!pip install scikitplot
```

### 0.1.3 Methodology

1. Explain your big data methodology
2. Introduce the topics you used in your project
  - Model 1
    - KNN
  - Model 2
    - Linear Regression

#### 0.1.4 Your code starts here

```
[58]: import pandas as pd, numpy as np, matplotlib.pyplot as plt, seaborn as sns
      from sklearn.neighbors import KNeighborsRegressor
      from sklearn.linear_model import LinearRegression
      from sklearn.model_selection import train_test_split, cross_val_score
      from sklearn.preprocessing import StandardScaler
      from sklearn.metrics import r2_score, mean_squared_error

      import warnings
      warnings.filterwarnings("ignore")
      %matplotlib inline
      sns.set_style("whitegrid")
```

```
[4]: df = pd.read_csv("data/house_data.csv")
```

```
[5]: # we're immediately just going to scrap any columns with missing values
      df = df.dropna(0)

      # clean up date

      df["date"] = pd.to_datetime(df["date"])
      df["year"] = df["date"].dt.year
      df["month"] = df["date"].dt.month

      # Split into X & y - we need to drop id, date, price, zipcode from X
      # date is now year/month, id is irrelevant, and i chose to not include the
      ↪ categorical data (zipcodes)
      X, y = df, df["price"]
      X = X.drop(["id", "date", "price", "zipcode"], axis=1)
```

```
[44]: # split dataset into train/test, proportioned at 25%
      scaler = StandardScaler()
      scaled_X = scaler.fit_transform(X)

      X_train, X_test, y_train, y_test = train_test_split(scaled_X, y, shuffle =
      ↪ True, test_size = 0.25)
```

```
[7]: df.head()
```

```
[7]:      id      date      price  bedrooms  bathrooms  sqft_living  \
0  7129300520  2014-10-13  221900.0         3         1.00         1180
1  6414100192  2014-12-09  538000.0         3         2.25         2570
2  5631500400  2015-02-25  180000.0         2         1.00          770
3  2487200875  2014-12-09  604000.0         4         3.00        1960
4  1954400510  2015-02-18  510000.0         3         2.00        1680
```

```
sqft_lot  floors  waterfront  view  ...  sqft_basement  yr_built  \
```

|   |       |     |   |   |     |     |      |
|---|-------|-----|---|---|-----|-----|------|
| 0 | 5650  | 1.0 | 0 | 0 | ... | 0   | 1955 |
| 1 | 7242  | 2.0 | 0 | 0 | ... | 400 | 1951 |
| 2 | 10000 | 1.0 | 0 | 0 | ... | 0   | 1933 |
| 3 | 5000  | 1.0 | 0 | 0 | ... | 910 | 1965 |
| 4 | 8080  | 1.0 | 0 | 0 | ... | 0   | 1987 |

|   | yr_renovated | zipcode | lat     | long     | sqft_living15 | sqft_lot15 | year | \ |
|---|--------------|---------|---------|----------|---------------|------------|------|---|
| 0 | 0            | 98178   | 47.5112 | -122.257 | 1340          | 5650       | 2014 |   |
| 1 | 1991         | 98125   | 47.7210 | -122.319 | 1690          | 7639       | 2014 |   |
| 2 | 0            | 98028   | 47.7379 | -122.233 | 2720          | 8062       | 2015 |   |
| 3 | 0            | 98136   | 47.5208 | -122.393 | 1360          | 5000       | 2014 |   |
| 4 | 0            | 98074   | 47.6168 | -122.045 | 1800          | 7503       | 2015 |   |

|   | month |
|---|-------|
| 0 | 10    |
| 1 | 12    |
| 2 | 2     |
| 3 | 12    |
| 4 | 2     |

[5 rows x 23 columns]

```
[8]: df.describe()
```

```
[8]:
```

|       | id           | price        | bedrooms     | bathrooms    | sqft_living  | \ |
|-------|--------------|--------------|--------------|--------------|--------------|---|
| count | 2.161300e+04 | 2.161300e+04 | 21613.000000 | 21613.000000 | 21613.000000 |   |
| mean  | 4.580302e+09 | 5.400881e+05 | 3.370842     | 2.114757     | 2079.899736  |   |
| std   | 2.876566e+09 | 3.671272e+05 | 0.930062     | 0.770163     | 918.440897   |   |
| min   | 1.000102e+06 | 7.500000e+04 | 0.000000     | 0.000000     | 290.000000   |   |
| 25%   | 2.123049e+09 | 3.219500e+05 | 3.000000     | 1.750000     | 1427.000000  |   |
| 50%   | 3.904930e+09 | 4.500000e+05 | 3.000000     | 2.250000     | 1910.000000  |   |
| 75%   | 7.308900e+09 | 6.450000e+05 | 4.000000     | 2.500000     | 2550.000000  |   |
| max   | 9.900000e+09 | 7.700000e+06 | 33.000000    | 8.000000     | 13540.000000 |   |

|       | sqft_lot     | floors       | waterfront   | view         | condition    | \ |
|-------|--------------|--------------|--------------|--------------|--------------|---|
| count | 2.161300e+04 | 21613.000000 | 21613.000000 | 21613.000000 | 21613.000000 |   |
| mean  | 1.510697e+04 | 1.494309     | 0.007542     | 0.234303     | 3.409430     |   |
| std   | 4.142051e+04 | 0.539989     | 0.086517     | 0.766318     | 0.650743     |   |
| min   | 5.200000e+02 | 1.000000     | 0.000000     | 0.000000     | 1.000000     |   |
| 25%   | 5.040000e+03 | 1.000000     | 0.000000     | 0.000000     | 3.000000     |   |
| 50%   | 7.618000e+03 | 1.500000     | 0.000000     | 0.000000     | 3.000000     |   |
| 75%   | 1.068800e+04 | 2.000000     | 0.000000     | 0.000000     | 4.000000     |   |
| max   | 1.651359e+06 | 3.500000     | 1.000000     | 4.000000     | 5.000000     |   |

|       | ... | sqft_basement | yr_built     | yr_renovated | zipcode      | \ |
|-------|-----|---------------|--------------|--------------|--------------|---|
| count | ... | 21613.000000  | 21613.000000 | 21613.000000 | 21613.000000 |   |
| mean  | ... | 291.509045    | 1971.005136  | 84.402258    | 98077.939805 |   |


|     |     |             |             |             |              |
|-----|-----|-------------|-------------|-------------|--------------|
| std | ... | 442.575043  | 29.373411   | 401.679240  | 53.505026    |
| min | ... | 0.000000    | 1900.000000 | 0.000000    | 98001.000000 |
| 25% | ... | 0.000000    | 1951.000000 | 0.000000    | 98033.000000 |
| 50% | ... | 0.000000    | 1975.000000 | 0.000000    | 98065.000000 |
| 75% | ... | 560.000000  | 1997.000000 | 0.000000    | 98118.000000 |
| max | ... | 4820.000000 | 2015.000000 | 2015.000000 | 98199.000000 |

|       | lat          | long         | sqft_living15 | sqft_lot15    | year \       |
|-------|--------------|--------------|---------------|---------------|--------------|
| count | 21613.000000 | 21613.000000 | 21613.000000  | 21613.000000  | 21613.000000 |
| mean  | 47.560053    | -122.213896  | 1986.552492   | 12768.455652  | 2014.322954  |
| std   | 0.138564     | 0.140828     | 685.391304    | 27304.179631  | 0.467616     |
| min   | 47.155900    | -122.519000  | 399.000000    | 651.000000    | 2014.000000  |
| 25%   | 47.471000    | -122.328000  | 1490.000000   | 5100.000000   | 2014.000000  |
| 50%   | 47.571800    | -122.230000  | 1840.000000   | 7620.000000   | 2014.000000  |
| 75%   | 47.678000    | -122.125000  | 2360.000000   | 10083.000000  | 2015.000000  |
| max   | 47.777600    | -121.315000  | 6210.000000   | 871200.000000 | 2015.000000  |

|       | month        |
|-------|--------------|
| count | 21613.000000 |
| mean  | 6.574423     |
| std   | 3.115308     |
| min   | 1.000000     |
| 25%   | 4.000000     |
| 50%   | 6.000000     |
| 75%   | 9.000000     |
| max   | 12.000000    |

[8 rows x 22 columns]

```
[13]: plt.figure(figsize = (16,16))
sns.heatmap(df.iloc[:, :-1].corr(), annot=True)
plt.show()
```



```
[41]: fig = plt.figure(figsize=(16, 16))
labels = list(df.drop(["id", "date", "view", "waterfront", "year", "month", "zipcode"], axis=1).columns)
for label in range(len(labels)):
    ax = fig.add_subplot(4,4,label+1)
    ax.plot(df.iloc[:, label], color = "blue", alpha = 0.3)
    ax.set_xlabel(labels[label])
    ax.autoscale(True)
```



```
[ ]: # Here is where the ML part starts
```

```
[45]: X_train
```

```
[45]: array([[ 0.67648506,  0.1756067 , -0.04344391, ..., -0.04645754,
          1.44790136, -1.46840343],
          [-0.39873715, -0.14900736, -0.23943274, ..., -0.27917455,
          1.44790136, -1.14740043],
          [ 1.75170727,  0.1756067 , -0.02166737, ..., -0.29184689,
          -0.69065478,  0.77861755],
          ...,
          [-0.39873715,  0.50022075, -0.65318696, ..., -0.42087774,
```

```

-0.69065478, -0.18439144],
[-1.47395936, -1.44746357, -1.39358923, ..., -0.27075073,
-0.69065478, 1.42062355],
[-1.47395936, -1.44746357, -0.78384618, ..., -0.00616973,
-0.69065478, 0.13661156]])

```

```
[46]: X_test
```

```

[46]: array([[ -1.47395936, -1.44746357, -0.37009197, ..., -0.17098348,
-0.69065478, 1.09962055],
[ -0.39873715, 0.8248348 , 0.66647122, ..., -0.25888414,
1.44790136, -1.14740043],
[ -1.47395936, -1.44746357, -1.01249983, ..., -0.07612401,
-0.69065478, 1.09962055],
...,
[ -1.47395936, -1.44746357, -1.28470655, ..., -0.27778279,
1.44790136, -1.78940643],
[ -0.39873715, 0.1756067 , -0.30476236, ..., -0.15190171,
-0.69065478, 1.09962055],
[ 0.67648506, -0.14900736, -0.1849914 , ..., 0.1037061 ,
-0.69065478, 0.45761455]])

```

```

[47]: K = []
scores = []

for i in range(2, 10):
    knn = KNeighborsRegressor(n_neighbors = i)
    knn.fit(X_train, y_train)
    s = knn.score(X_test, y_test)
    print(f"k = {i}: {s}")
    scores.append(s)

print(scores)

```

```

k = 2: 0.7414030057849562
k = 3: 0.766575752224133
k = 4: 0.7742231283592205
k = 5: 0.7812822298714657
k = 6: 0.7844945931859527
k = 7: 0.7878581007483119
k = 8: 0.7894831792276535
k = 9: 0.7904318377078658
[0.7414030057849562, 0.766575752224133, 0.7742231283592205, 0.7812822298714657,
0.7844945931859527, 0.7878581007483119, 0.7894831792276535, 0.7904318377078658]

```

```

[48]: good_k = np.argmax(scores) + 1

knn = KNeighborsRegressor(n_neighbors = good_k)

```

```
knn.fit(X_train, y_train)
```

```
[48]: KNeighborsRegressor(n_neighbors=8)
```

```
[49]: y_pred = knn.predict(X_test)
```

```
[50]: r2_score(y_test, y_pred)
```


```
[50]: 0.7894831792276535
```

```
[51]: mean_squared_error(y_test, y_pred)
```

```
[51]: 25681060979.49989
```

```
[57]: sns.scatterplot(x = y_test, y = y_pred)  
plt.xlabel("actual price")  
plt.ylabel("actual price")
```

```
[57]: Text(0, 0.5, 'actual price')
```





```
[67]: reg = LinearRegression()
      reg.fit(X_train, y_train)

      reg.score(X_test, y_test), reg.coef_
```

```
[67]: (0.6925495620411505,
      array([-39241.12953713,  34324.78619446,  87410.07508251,   6321.57012566,
            -2571.62511083,  47898.12687306,  38942.01578742,  20425.65307168,
            113778.66497505,  82184.29435247,  27622.24240613, -70742.14923179,
             9304.95451654,  76900.85415064, -18174.16249917,  11997.57351894,
            -10604.10639166,  17626.82769587,   4711.39783398]))
```

```
[63]: y_pred = reg.predict(X_test)
```

```
[64]: r2_score(y_test, y_pred)
```


```
[64]: 0.6925495620411505
```

```
[65]: mean_squared_error(y_test, y_pred)
```

```
[65]: 37506045438.21011
```

```
[66]: sns.scatterplot(x = y_test, y = y_pred)
      plt.xlabel("actual price")
      plt.ylabel("actual price")
```

```
[66]: Text(0, 0.5, 'actual price')
```



### 0.1.5 Conclusions

```
[71]: print("""
Overall, KNN was obviously much better - I'm not really surprised. It's
    definitely more involved than plain linear regression.

The r2 score was lower for both than I was hoping for - loosely, it means that
    the dataset accounted for 78% of the variance (in the KNN regressor), and
    then only 69% in the linear regression. This is probably not entirely
    accurate, since I'd say house prices are close to 90% at minimum, but who
    knows - maybe I know something the real estate market doesn't!

The MSE is abysmal because house prices are so high. Unfortunately, it's not a
    very useful metric here - however, it would be remiss to NOT show it,
    because it's a staple metric of regression problems.""")
```

Overall, KNN was obviously much better - I'm not really surprised. It's definitely more involved than plain linear regression.

The  $r^2$  score was lower for both than I was hoping for - loosely, it means that the dataset accounted for 78% of the variance (in the KNN regressor), and then only 69% in the linear regression. This is probably not entirely accurate, since I'd say house prices are close to 90% at minimum, but who knows - maybe I know something the real estate market doesn't!

The MSE is abysmal because house prices are so high. Unfortunately, it's not a very useful metric here - however, it would be remiss to NOT show it, because it's a staple metric of regression problems.

### 0.1.6 References

- Academic (if any)
- Online (if any)

```
[69]: print("""None applicable, I think.""")
```

None applicable, I think.

### 0.1.7 Credits

- If you use and/or adapt your code from existing projects, you must provide links and acknowledge the authors. > *This code is based on ... (if any)*

```
[68]: print("""This code is based on the template that existed here before me, and
↳also the sklearn docs (that I had to double check a bunch of the exact
↳packages things were in). I used no other external sources for code help.""")
```

This code is based on the template that existed here before me, and also the sklearn docs (that I had to double check a bunch of the exact packages things were in). I used no other external sources for code help.

```
[ ]: # End of Project
```